

1 Math Fundamentals

Formulae

Asymptotics Let f and g be non-negative functions, then $f(n)$ is...

	...if and only if...	if $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)}$
$O(g(n))$	$\exists c, N : f(n) \leq c \cdot g(n) \text{ for } n \geq N$	$\neq \infty$
$o(g(n))$	$\forall c, \exists N : f(n) \leq c \cdot g(n) \text{ for } n > N$	$= 0$
$\Omega(g(n))$	$\exists c, N : f(n) \geq c \cdot g(n) \text{ for } n \geq N$	$\neq 0$
$\omega(g(n))$	$\forall c, \exists N : f(n) \geq c \cdot g(n) \text{ for } n > N$	$= \infty$
$\Theta(g(n))$	$f(n)$ is $O(g(n))$ and $g(n)$ is $O(f(n))$	$\neq 0, \infty$

Implications	Equivalences
$f = o(g) \Rightarrow f = O(g)$	$f = O(g) \Leftrightarrow g = \Omega(f)$
$f = \omega(g) \Rightarrow f = \Omega(g)$	$f = o(g) \Leftrightarrow g = \omega(f)$
	$f = O(g) \text{ and } f = \Omega(g) \Leftrightarrow f = \Theta(g)$

Generalizations: For all $a, b > 0, k > 1$, and $n \geq 1$

$$\begin{cases} (\log n)^b = o(n^a) \\ n^b = o((1+a)^n) \\ a^{\sqrt{\log n}} = o(n^b) \\ k^n > n! > n^{n^b} > n^a > \log n > n^{1/k} > O(1) \end{cases} \approx \text{comparisons}$$

Stirling's Formula

$$n! \approx n^n \cdot e^{-n} \cdot \sqrt{2\pi n}$$

Master Theorem: Let $T(n) = aT(n/b) + cn^k$ for $a \geq 1, b \geq 2, c, k \geq 0$. Then $T(n)$ is

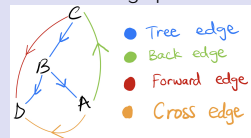
$$\begin{cases} \Theta(n^k) & \text{if } a < b^k \\ \Theta(n^k \log n) & \text{if } a = b^k \\ \Theta(n^{\log_b a}) & \text{if } a > b^k \end{cases}$$

2 Graph Algorithms

Depth First Search Algorithm

DFS Algorithm **Runtime:** $O(|V| + |E|)$

- Goes as deep as possible in the graph then backtrack.
- (In-class algorithm) marks preorder and postorder numbers of each node.
- In a directed graphs, DFS will label edges as follows. No cross edges in undirected graphs.



Applications:

Detecting cycles
Topological Sort on **DAGs**
Strongly Connected Components

$\exists (u, v): \text{postorder}(u) < \text{postorder}(v)$
Decreasing post-order
Hint: DFS on reversed Graph.

Single Source Shortest Paths

BFS Algorithm **only for unweighted graphs.**

- Goes level by level in the graph.
- Uses a queue to process nodes.

Runtime: $O(|V| + |E|)$

Let **update(u, v)** be defined as: if $\text{Dist}[u] + \text{length}(u, v) < \text{Dist}[v]$, update: $\text{Prev}[v] = u$ and $\text{Dist}[v] = \text{Dist}[u] + \text{length}(u, v)$.

Dijkstra's Algorithm: **only +ve weighted graphs**

- Let $\text{Dist}[v] = \infty$ and $\text{Prev}[v] = \text{NIL}$ for all vertices v .
- Let $\text{Dist}[s] = 0$ and initialize MinHeap with $(s, 0)$.
- While heap isn't empty, keep popping the vertex (say u) with the smallest distance. For each neighbor v of u with weight w , **update(u, v)**.

Runtime: $O(|V| * \text{popMin} + |E| * \text{Insert})$

Bellman-Ford Algorithm: **Only + and -ve weighted graphs**

- Let $\text{Dist}[v] = \infty$ and $\text{Prev}[v] = \text{NIL}$ for all vertices v .
- Let $\text{Dist}[s] = 0$.
- For $|V| - 1$ times: For each edge (u, v) , **update(u, v)**
- Checking for negative weight cycles: repeat the above step once. If any distance can still be improved, a negative weight cycle exists. Hence, return **Inconclusive**

Runtime: $O(|V| * |E|)$

Linear Runtime: **only for Directed Acyclic Graphs (DAGs)**

- Run DFS to get topological sort.
- For every edge (u, v) , in topological sort, **update(u, v)**.

Runtime: $O(|V| + |E|)$

Min Heaps

Representation: Visually a complete tree. Implementationwise, let $A[0..n-1]$ be a list where $A[0]$ is the root and for any i th element, its parent, left, and right children are at $\lfloor i/2 \rfloor, 2i, 2i+1$ respectively.

Heap Property: The parent element is smaller than its children.

Operations	Description
$\text{Insert}(a)$	let $A[n] = a$ and $\text{HeapifyUp}(n)$
PopMin	let $A[0] = A[n-1]$ and $\text{HeapifyDown}(0)$
$\text{HeapifyUp}(i)$	repeatedly swap $A[i]$ with its parent until the heap property is restored
HeapifyDown	repeatedly swap $A[i]$ with its smallest child until the heap property is restored

Heaps Operations and Runtimes: Both $O(\log n)$ with binary heaps.

Note: Can do better with Fibo-Heaps (amortized $O(1)$ for PopMin.)

Minimum Spanning Trees

Basic Properties:

- a **Tree** is connected, acyclic, and has $|V| - 1$ edges (any two implies the third).
- Cut Property** states that for any cut of a connected, undirected graph, the minimum weight edge that crosses the cut belongs to the MST.
- Only for connected, undirected, and weighted (non-negative) graphs.

Prim's Algorithm:

- Start with a single vertex and greedily add closest vertices.
- Similar to Dijkstra's algorithm, but $\text{dist}[v]$ is the weight of the edge connecting v to the MST instead of the distance from s .

Runtime: $O(|E| \log |V|)$ with Fibo-heaps

Kruskal's Algorithm:

- Sort edges in ascending order of weight.
- Repeatedly add the lightest edge that does not create a cycle until we have $|V| - 1$ edges.

Runtime: $nT(\text{Union}) + mT(\text{Find}) + T(\text{Sort } m \text{ Edges})$.

Notes: Implemented using a union-find data structure.

Disjoint Forest Data Structure

Maintain disjoint sets that can be combined ("unioned") efficiently. Operations $\text{MakeSet}(x)$, $\text{Find}(x)$, and $\text{Union}(a, b)$

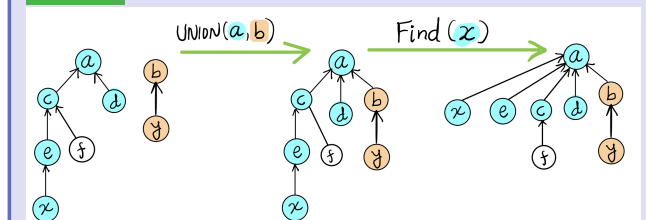
Runtime: Any sequence of m UNIONS and n FINDs operations take $O((m+n) \log^* n)$.

Note: $\log^* n$ is the number of times we can $\log_2 n$ until we get ≤ 1 .

Heuristics:

- Union by rank. When performing a union operation, we prefer to merge the shallower tree into the deeper tree.
- Path compression. After performing a find operation, attach all the nodes touched directly onto the root of the tree.

Example:



3 Greedy Algorithms

Main idea: At each step, make a locally optimal choice in hope of reaching the globally optimal solution.

Horn Formula

Algorithm: Set all variables to false and greedily set ones to be true when forced to.

Runtime: linear time in the length of the formula (i.e., the total number of appearances of all literals).


Notes: Only works for SAT instances where in each clause, there is at most one positive literal.

Huffman Coding

Algorithm: **Runtime:** $O(n \log n)$

Find the best encoding by greedily combining the two least frequently items. Optimal in terms of encoding one character at a time.

Example: A Huffman tree for string "Mississippi River"



Set Cover

Given $X = \{x_1, \dots, x_n\}$, and a collection of subsets \mathcal{S} of X such that $\bigcup_{S \in \mathcal{S}} S = X$, find the subcollection $\mathcal{T} \subseteq \mathcal{S}$ such that the sets of \mathcal{T} cover X .

Algorithm: **Runtime:** $O(|U|)$

1. Greedily choose the set that covers the most number of the remaining uncovered elements at the given iteration.

claim: Let k be the size of the smallest set cover for the instance (X, \mathcal{S}) . Then the greedy heuristic finds a set cover of size at most $k \ln n$.

Note: Not always optimal; achieves $O(\log n)$ approximation ratio.

4 Divide and Combine

Main Idea: Divide the problem into smaller pieces, recursively solve those, and then combine their results to get the final result.

Famous Examples w/ Runtimes

Mergesort	$O(n \log n)$
Min and Max on a line	$\frac{3}{2}n - 2$ comparisons; $O(n)$ runtime.
Closest Pair of Points	$O(n \log^2 n)$

n -digit Integer Multiplication

standard Multiplication	$\Theta(n^2)$
3 products on $n/2$ digits	$\Theta(n^{\log_2 3}) = \Theta(n^{1.59})$
5 products on $n/3$ digits	$\Theta(n^{\log_3 5}) = \Theta(n^{1.46})$

$n \times n$ Matrix Multiplication

Strassen's Algorithm: **Runtime:** $O(n^{\log_2 7})$

Divide into four submatrices, each of size $n/2$ by $n/2$.

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix} \begin{bmatrix} E & F \\ G & H \end{bmatrix} = \begin{bmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{bmatrix}$$

Find: $P_1 = A(F - H)$, $P_2 = (A + B)H$, $P_3 = (C + D)E$, $P_4 = D(G - E)$, $P_5 = (A + D)(E + H)$, $P_6 = (B - D)(G + H)$, $P_7 = (C - A)(E + F)$, then:

$AE + BG = -P_2 + P_4 + P_5 + P_6$ and $AF + BH = P_1 + P_2$
 $CE + DG = P_3 + P_4$ and $CF + DH = P_1 - P_3 + P_5 + P_7$

5 Dynamic Programming

Main Idea: Maintain a lookup table of correct solutions to subproblems and build up this table towards the actual solution.

Steps:

1. Define subproblems and recurrence to solve subproblems.
2. Combine with **reuse**.
3. Runtime and space analysis.

Edit Distance

Find the minimum number of operations required to transform one string, $A[1 \dots n]$, into another, $B[1 \dots m]$.

Algorithm: **Runtime and Space:** $O(nm)$

1. Subproblem: let $D(i, j)$ represent the edit distance between $A[1 \dots i]$ and $B[1 \dots j]$.
2. Recurrence is:
Base cases: $D(i, 0) = i$, $D(0, j) = j$.
 $D(i, j) = \min[D(i-1, j) + 1, D(i, j-1) + 1, D(i-1, j-1) + (1 \text{ if } i=j, 0 \text{ otherwise})]$.
3. return $D(n, m)$.

All Pairs Shortest Paths

Given a graph G with n vertices and m edges, calculate distances of the shortest paths between every pair of nodes.

Floyd-Warshall Algorithm: **Runtime:** $O(n^3)$

1. Subproblem: let $D_k[i, j]$ represent the shortest path between i and j using only nodes in $[1 \dots k]$.
2. Recurrence is:
$$\begin{cases} D_0[i, j] = d_{ij} \text{ if } i \text{ and } j \text{ are connected, } \infty \text{ otherwise.} \\ D_k[i, j] = \min(D_{k-1}[i, j], D_{k-1}(i, k) + D_{k-1}[k, j]). \end{cases}$$
3. return $D(i, j, n)$.

Notes: Does not work for cyclic graphs.

Hashing and Set Resemblance

Primality Testing

Algorithm: Generate large (d-digit) primes: Generate a random d-digit number. Check if it is prime. If not, repeat.

Facts 1: k^{th} prime number is $\Theta(k \log k)$.

Fact 2: of the integers $1, \dots, n$, $\Theta(n/\log n)$ are prime.

How many d-digits generations until prime: $O(d)$.

Fermat's Little Theorem:

If p is prime and a is not divisible by p , then $a \in \mathbb{Z}$, $a^{p-1} \equiv 1 \pmod{p}$.

Notes: a -Pseudoprime p : $a^{p-1} \equiv 1 \pmod{p}$, but p is not prime.

Carmichael number: composite number n s.t. $a^{n-1} \equiv 1 \pmod{n}$ for all $a \in \mathbb{Z}$. Example: 561

Miller-Rabin Primality Test:

If p is prime, the only solutions to $a^2 \equiv 1 \pmod{p}$ are $a \in \{\pm 1\}$.

A **non-trivial square root of 1** is an integer a such that $a^2 \equiv 1 \pmod{n}$ and $a \not\equiv \pm 1 \pmod{n}$.

The Test

1. Choose a random $a \in [n]$
2. if $a^{n-1} \not\equiv 1 \pmod{n}$, return n is composite and a is a witness.
3. Let $n-1 = 2^t u$ and compute $a^u, a^{2u}, \dots, a^{2^{t-1}u}$.
4. Check for $a^{2^{i-1}u} \not\equiv \pm 1 \pmod{n}$, $a^{2^i u} \equiv 1 \pmod{n}$
5. If so, we've found a *non-trivial square root of 1* modulo n .

a is a **witness** to the compositeness of n , if n fails the test under a .

Note: Let \mathbf{F} be *incorrectly* identifying n as prime. $\Pr(\mathbf{F}) \leq 1/4$. With k independent tests, $\Pr(\mathbf{F}) \leq 4^{-k}$.

Euclid's algorithm

Def: $\gcd(a, b) = \max\{d \in \mathbb{Z} : d|a \text{ and } d|b\}$.

Basic algorithm: find the $\gcd(a, b)$ by repeatedly subtracting the smaller number from the larger number.

Extended Euclid's algorithm: find x, y such that $ax + by = \gcd(a, b)$.

```
def gcd(a, b):      # a >= b >= 0
    if b == 0:
        return a
    return gcd(b, a mod b)
```

```
def Extended-Euclid(a, b):      # a >= b >= 0
    if b == 0: return (a, 1, 0)
    Compute k such that a = bk + (a mod b)
    (d, x, y) = Extended-Euclid(b, a mod b)
    return (d, y, x-ky)
```

Example: Run extended algorithm on $a = 35, b = 15$.

- $\gcd(35, 15) = \gcd(15, 5) = \gcd(5, 0) = 5$.
- $35 = 2 \cdot 15 + 5$.
- $15 = 3 \cdot 5 + 0$.
- $5 = 1 \cdot 0 + 5$.
- $x = 1, y = -2$.

NP-c approximations

α -approximation: f is an α -approximation to f^* if $f(x) \leq \alpha f^*(x)$ for all x .

Independent Set: Given a graph $G = (V, E)$, find largest set of vertices $S \subseteq V$ such that no two vertices in S are adjacent.

Vertex Cover: Given a graph $G = (V, E)$, find the smallest set of vertices $S \subseteq V$ such that every edge in E is incident to at least one vertex in S .

Max Cut: Given a graph, divide vertices into two sets to maximize number of edges between them.

- Split vertices arbitrarily.
- While moving a vertex improves the solution, move it.
- stop when no more moves improve the solution.

NP-c Heuristics

MAX SAT: Linear Relaxation

- Convert formula to integer equations. E.g. $(X \vee \bar{Y}) \rightarrow x' + (1 - y') \geq 1$.
- Relax the constraint for variables $x', y', \dots \in \{0, 1\}$ to $x', y', \dots \in [0, 1]$.
- Solve the relaxed problem to get assignment x', y', \dots . Then let $X = 1$ with probability x'

Claim: If formula is satisfiable, then the number of clauses satisfied by above algorithm $\geq |\text{clauses}| \cdot (1 - 1/e) \approx 0.63 |\text{clauses}|$

MAX SAT: Local Search

- Pick a random assignment $x, y, \dots \in \{0, 1\}$.
- While moving a variable improves the solution, move it.
- stop when no more moves improve the solution.

Cryptography

Receiver: picks two large prime numbers p and q , compute $n = pq$.

- public key: (n, e) where e is random s.t. $\gcd(e, (p-1)(q-1)) = 1$.
- private key: $d = e^{-1} \pmod{(p-1)(q-1)} \implies de \equiv 1 \pmod{(p-1)(q-1)}$.

RSA algorithm: Let sender's message be m .

1. Encryption: $y = E(m) = m^e \pmod{n}$ —using public key (n, e)
2. Decryption: $m = D(y) = y^d \pmod{n}$ —using private key d

Linear Programming

Simplex algorithm: Hill climbing on feasible region always yields an optimal solution.

Network Flow

Story: Given the number of available tickets t_{ij} between cities i and j , along with the city map, the goal is to maximize the number of tickets sold for people traveling from city s to city t .

Reduction to LP: Objective function:

Maximize the total flow of tickets from the source node s to the target node t . In other words, maximize the sum of x_{st} over all edges (s, t) .

$$\text{maximize } \sum_{(s,t)} x_{st} \quad (1)$$

Subject to the following constraints:

1. Flow conservation constraints: For each node i , other than s and t , the flow into the node must equal the flow out of the node.

$$\sum_j x_{ij} - \sum_k x_{ki} = 0, \forall i \neq s, t \quad (2)$$

2. Capacity constraints: The flow of tickets on each edge (i, j) must not exceed the available tickets t_{ij} .

$$0 \leq x_{ij} \leq t_{ij}, \forall (i, j) \quad (3)$$

Ford Fulkerson:

- Make a Residual graph and Start with empty flow.
- While there is an augmenting path from s to t :
 - Find the bottleneck capacity b on the path.
 - Increase the flow on the path by b and update residues.
- The final flow is the maximum flow.

Augmenting path: is a path from s to t such that the flow on the path can be increased by at least one unit.

Residual Graph: for every edge (x,y) of G , add edge (y,x) , capacity $c(y,x)$ = flow $f(x,y)$

Runtime:

- with DFS to find augmenting paths $O(V EU)$, where U is max edge capacity.
- $O(V E^2)$ with BFS ("Edmonds-Karp algorithm").

Max-flow min-cut theorem: The maximum flow is equal to the minimum capacity of an s - t cut.

Claim: when Ford-Fulkerson terminates, we can find a cut matching the flow.

- No s - t path (of nonzero edges) in residual graph.
- Choose V_1 = vertices reachable from s
- All edges e leaving V_1 have $f(e) = c(e)$
- All edges e entering V_1 have $f(e) = 0$
- So, total flow = capacity of cut.